

Regression

4. Batch non-linear projection methods

Olivier Sigaud


Sorbonne Université
<http://people.isir.upmc.fr/sigaud>



Radial Basis Function Networks



Reminder: Outline of methods




- ▶ Projecting the input space into a feature space using non-linear basis functions (shown with RBFNs)



Stulp, F. and Sigaud, O. (2015). Many regression algorithms, one unified model: A review. *Neural Networks*, 69:60–79.


Least Square Projection Methods: framework



- ▶ With linear regression, we look for $\hat{f}(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b$
- ▶ This is not general enough for non-linear functions
- ▶ More general form: $\hat{f}(\mathbf{x}) = \sum_{e=0}^E w_e \cdot \phi_{\theta_e}(\mathbf{x})$ with $\phi_{\theta_0}(\mathbf{x}) = 1$
- ▶ This can be seen as projecting the input to a different space...
- ▶ ... where the latent function is linear



Understanding projection




- ▶ Given the functions $\phi_{\theta_e}(\cdot)$, each point \mathbf{x} in the dataset \mathbf{X} is projected into a point $\mathbf{x}' = \phi_{\theta_e}(\mathbf{x})$
- ▶ The number E of functions determines the dimension of \mathbf{x}'
- ▶ The point $\mathbf{x} = (x_1, x_2)$ is projected to $\mathbf{x}' = (\phi_{\theta_1}(\mathbf{x}), \phi_{\theta_2}(\mathbf{x}), \phi_{\theta_3}(\mathbf{x}))$




Cybenko, G. (1989) Approximation by superpositions of a sigmoidal function. *Mathematics of Control, Signals, and Systems (MCSS)*, 2(4):303–314

Basis Function Networks: finding the weights



- ▶ Now we have a dataset \mathbf{X}', \mathbf{y}
- ▶ We want to find the function $y = \hat{f}(\mathbf{x}') = \sum_{e=0}^E w_e \cdot \mathbf{x}'$
- ▶ This is a linear regression problem
- ▶ Thus we perform regression in the projected space
- ▶ The larger E , the smaller the error (perfect fit in the limit of $E \rightarrow \infty$)

Standard features: Gaussian basis functions



- ▶ The more features, the better the approximation
- ▶ ... but the more expensive the computation
- ▶ If the features are given and constant, this is a linear architecture

Kernel Ridge Regression (KRR) = Kernel Regularised Least Squares (KRGLS)

- ▶ Define features with a kernel function $k(\mathbf{x}, \mathbf{x}_i)$ per point \mathbf{x}_i
- ▶ Define the Gram matrix as a kernel matrix:

$$\mathbf{K} = \begin{bmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & k(\mathbf{x}_1, \mathbf{x}_2) & \cdots & k(\mathbf{x}_1, \mathbf{x}_N) \\ k(\mathbf{x}_2, \mathbf{x}_1) & k(\mathbf{x}_2, \mathbf{x}_2) & \cdots & k(\mathbf{x}_2, \mathbf{x}_N) \\ \vdots & \vdots & \ddots & \vdots \\ k(\mathbf{x}_N, \mathbf{x}_1) & k(\mathbf{x}_N, \mathbf{x}_2) & \cdots & k(\mathbf{x}_N, \mathbf{x}_N) \end{bmatrix}. \quad (1)$$

- ▶ If we had an infinity of data points, the linear approximation in feature space would become perfect
- ▶ Intuition: the error is a function of the distance to data points
- ▶ Computing the weights is done with RR using

$$\boldsymbol{\theta}^* = (\lambda \mathbf{I} + \mathbf{K})^{-1} \mathbf{y}, \quad (2)$$

- ▶ Note that \mathbf{K} is symmetric
- ▶ The kernel matrix \mathbf{K} grows with the number of points (kernel expansion)
- ▶ The matrix inversion may become too expensive
- ▶ Solution: finite set of features (RBFNs), incremental methods

Radial Basis Function Networks: definition and solution

- ▶ Radial Basis Functions versus Kernels (Gaussians
 $\phi(\mathbf{x}, \boldsymbol{\theta}_e) = e^{-\frac{1}{2}(\mathbf{x}-\mu_e)^T \boldsymbol{\Sigma}_e^{-1} (\mathbf{x}-\mu_e)}$ are both)
- ▶ We define a set of E basis functions (often Gaussian)

$$\hat{f}(\mathbf{x}) = \sum_{e=1}^E w_e \cdot \phi(\mathbf{x}, \boldsymbol{\theta}_e) \quad (3)$$

$$= \boldsymbol{\theta}^\top \cdot \phi(\mathbf{x}). \quad (4)$$

- ▶ We also define the *Gram matrix*

$$\mathbf{G} = \begin{bmatrix} \phi(\mathbf{x}_1, \boldsymbol{\theta}_1) & \phi(\mathbf{x}_1, \boldsymbol{\theta}_2) & \cdots & \phi(\mathbf{x}_1, \boldsymbol{\theta}_E) \\ \phi(\mathbf{x}_2, \boldsymbol{\theta}_1) & \phi(\mathbf{x}_2, \boldsymbol{\theta}_2) & \cdots & \phi(\mathbf{x}_2, \boldsymbol{\theta}_E) \\ \vdots & \vdots & \ddots & \vdots \\ \phi(\mathbf{x}_N, \boldsymbol{\theta}_1) & \phi(\mathbf{x}_N, \boldsymbol{\theta}_2) & \cdots & \phi(\mathbf{x}_N, \boldsymbol{\theta}_E) \end{bmatrix} \quad (5)$$

- ▶ and we get the least squares solution

$$\boldsymbol{\theta}^* = (\mathbf{G}^\top \mathbf{G})^{-1} \mathbf{G}^\top \mathbf{y}.$$

Least Square Projection Methods: summary of computations

► Linear case

$$\boldsymbol{\theta}^* = (\bar{\mathbf{X}}^\top \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^\top \mathbf{y} \quad (LS) \quad (7)$$

$$\boldsymbol{\theta}^* = (\lambda \mathbf{I} + \bar{\mathbf{X}}^\top \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^\top \mathbf{y}. \quad (RLS) \quad (8)$$

► Gram matrix case

$$\boldsymbol{\theta}^* = (\mathbf{G}^\top \mathbf{G})^{-1} \mathbf{G}^\top \mathbf{y} \quad (RBFN) \quad (9)$$

► Kernel matrix case

$$\boldsymbol{\theta}^* = \mathbf{K}^{-1} \mathbf{y}, \quad (GPR) \quad (10)$$

$$\boldsymbol{\theta}^* = (\lambda \mathbf{I} + \mathbf{K})^{-1} \mathbf{y}. \quad (KRR) \quad (11)$$


LWR versus RBFNs

$$\hat{f}(\mathbf{x}) = \sum_{e=1}^E \phi(\mathbf{x}, \boldsymbol{\theta}_e) \cdot (b_e + \mathbf{a}_e^\top \mathbf{x}) \quad (12)$$

$$\hat{f}(\mathbf{x}) = \sum_{e=1}^E \phi(\mathbf{x}, \boldsymbol{\theta}_e) \cdot w_e, \quad (13)$$

- ▶ Eq. (13) is a special case of (12) with $\mathbf{a}_e = \mathbf{0}$ and $b_e = w_e$.
- ▶ RBFNs: performs one LS computation in a projected space
- ▶ LWR: performs many LS computation in local domains

Wrap-up



- Image taken from Freek Stulp's IROS 2018 Tutorial

Regression

- └ Radial Basis Function Networks
 - └ RBFNs



Bishop, C. M. (2007).

Pattern recognition and machine learning.

Springer Berlin/Heidelberg, Germany.



Stulp, F. and Sigaud, O. (2015).

Many regression algorithms, one unified model: A review.

Neural Networks, 69:60–79.