

Regression

5. Incremental and iterative methods

Olivier Sigaud

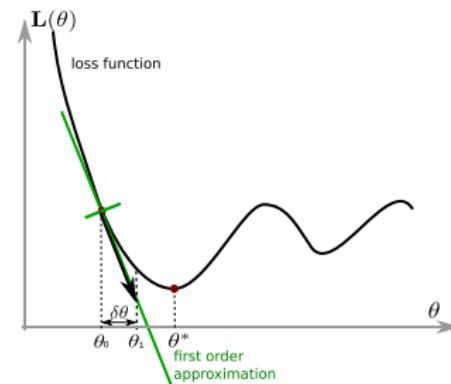
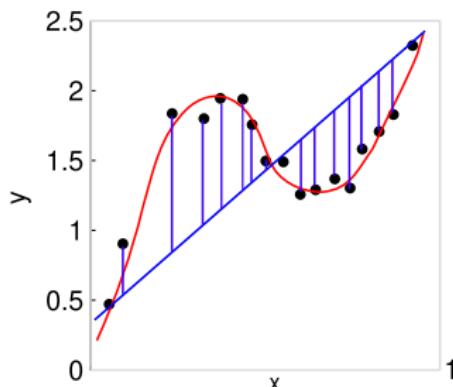
Sorbonne Université
<http://people.isir.upmc.fr/sigaud>

Gradient descent

Limiting the cost of regression

- ▶ Solving $\theta^* = (\mathbf{G}^\top \mathbf{G})^{-1} \mathbf{G}^\top \mathbf{y}$ (or any variant) requires inverting $(\mathbf{G}^\top \mathbf{G})$
- ▶ That is in $O(N^3)$ in the number of dimensions
- ▶ For large linear architectures or deep neural networks, this is too expensive
- ▶ Incremental solutions:
 - ▶ Complexity can be reduced to $O(N^2)$ by using the Sherman-Morrisson formula → incremental update of the inverse,
 - ▶ But sensitive to rounding errors.
 - ▶ Numerically more stable option: updating the Cholesky factor of the matrix using the QR algorithm.
 - ▶ In the linear case: Recursive Least Squares

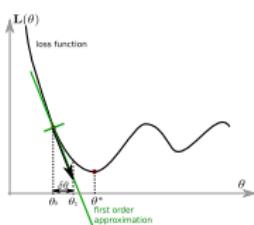
Regression through batch gradient descent



- ▶ We want to minimize a regression error (left)
- ▶ Thus find the minimum of a **loss function** $L(\theta)$ (right)
- ▶ Instead of solving $\theta^* = \min_{\theta} L(\theta)$ (in one step, but expensive)
- ▶ Rather perform local steps using an iterative approach
- ▶ The iteration equation is $\theta_{i+1} \leftarrow \theta_i + \delta\theta_i$
- ▶ This is the essence of **gradient descent**

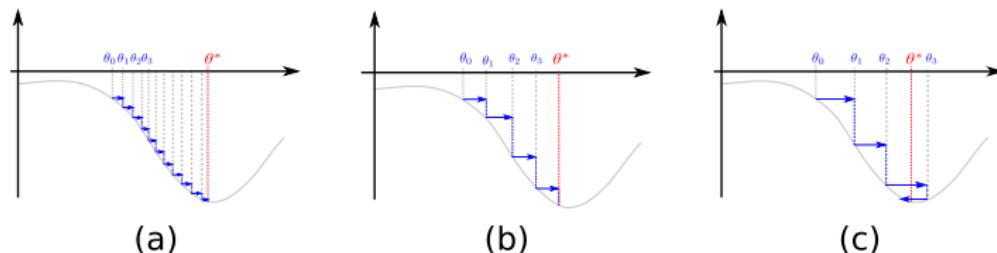
Gradient descent is steepest descent

- ▶ We want to minimize $\mathbf{L}(\boldsymbol{\theta}_i + \delta\boldsymbol{\theta}_i)$ over $\delta\boldsymbol{\theta}_i$
- ▶ How to choose $\delta\boldsymbol{\theta}_i$?



- ▶ The optimum of $\mathbf{J}(\boldsymbol{\theta} + \delta\boldsymbol{\theta})$ over $\delta\boldsymbol{\theta}$ is reached when $\frac{\partial \mathbf{J}(\boldsymbol{\theta} + \delta\boldsymbol{\theta})}{\partial \delta\boldsymbol{\theta}} = 0$
- ▶ First order approx: $\mathbf{J}(\boldsymbol{\theta}) + \nabla_{\boldsymbol{\theta}}\mathbf{J}(\boldsymbol{\theta})^T \delta\boldsymbol{\theta} + \nu \delta\boldsymbol{\theta}^T \delta\boldsymbol{\theta} + \text{higher order terms}$
- ▶ $\frac{\partial \mathbf{J}(\boldsymbol{\theta} + \delta\boldsymbol{\theta})}{\partial \delta\boldsymbol{\theta}} \sim \nabla_{\boldsymbol{\theta}}\mathbf{J}(\boldsymbol{\theta})^T \delta\boldsymbol{\theta} + 2\nu \delta\boldsymbol{\theta}$
- ▶ $\delta\boldsymbol{\theta}^* = -\alpha \nabla_{\boldsymbol{\theta}}\mathbf{J}(\boldsymbol{\theta})$
- ▶ Thus the steepest descent direction is given by the first order derivative

Tuning the step size

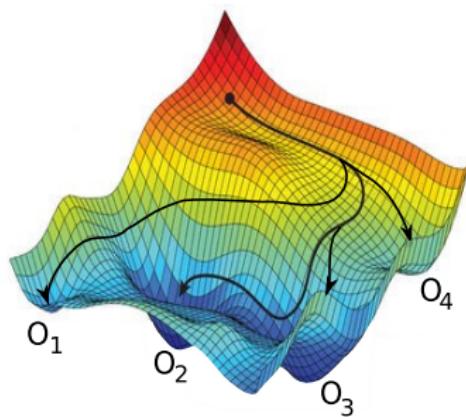


- ▶ Minimizing the first order derivative is not lower bounded
- ▶ We need a step size α to determine how far to go
- ▶ (a) α_i are too small, (b) α_i are adequate, (c) α_i are too large
- ▶ If too small, too many steps. If too large, may miss a local optimum
- ▶ Line search: iterate to find the best step size (used e.g. in TRPO)



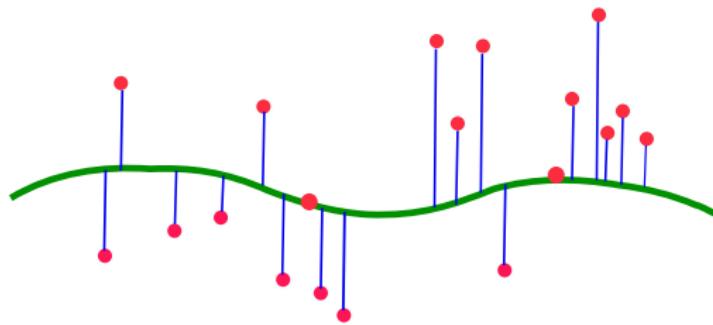
Schulman, J., Levine, S., Moritz, P., Jordan, M. I., & Abbeel, P. (2015) Trust region policy optimization. *CoRR*, [abs/1502.05477](https://arxiv.org/abs/1502.05477)

Local Optima



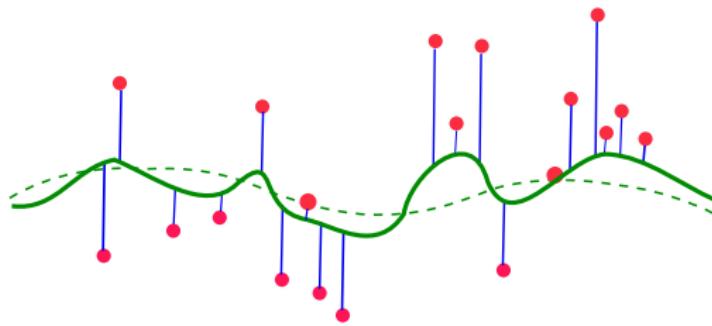
- ▶ Gradient descent is a local improvement approach where, at each step, we follow the steepest descent direction
- ▶ But we don't know where is the optimum we are targetting
- ▶ Unless the cost function is convex in the parameter space, gradient descent can end-up in different local optima depending on the starting point or noise
- ▶ Anything that basically goes down will end up somewhere low!

Regression through batch gradient descent: illustration



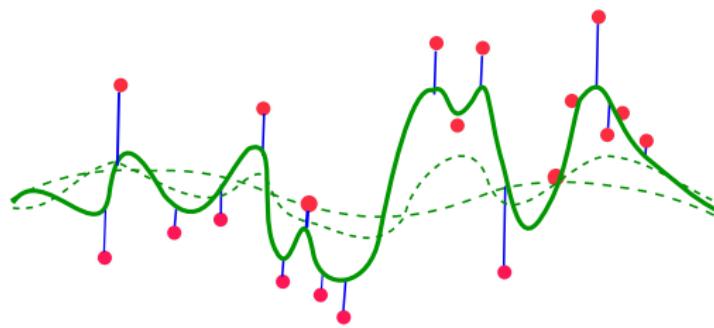
- ▶ As usual, the loss is the squared sum of all residual errors
$$L(\theta) = \|\mathbf{y} - \hat{f}(\mathbf{X})\|^2$$

Regression through batch gradient descent: illustration



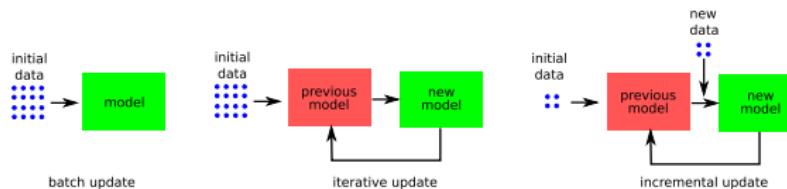
- ▶ As usual, the loss is the squared sum of all residual errors
$$L(\theta) = \|\mathbf{y} - \hat{f}(\mathbf{X})\|^2$$

Regression through batch gradient descent: illustration



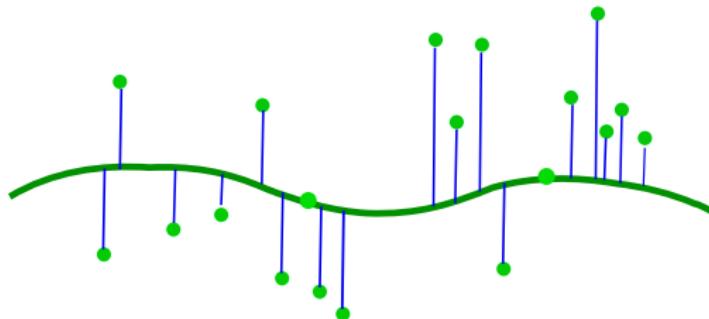
- ▶ As usual, the loss is the squared sum of all residual errors
$$L(\theta) = \|\mathbf{y} - \hat{f}(\mathbf{X})\|^2$$

Batch, Iterative, Incremental



- ▶ So far, we have focused on batch regression
- ▶ Iterative methods: you improve a model through steps, using the same data (your batch) at all steps
- ▶ Incremental methods: you improve a model through steps, with additional data at each step
- ▶ Incremental implies more than iterative
- ▶ Batch gradient descent is only iterative: no need for additional data
- ▶ Stochastic gradient descent is a specific case of incremental (y and X change for each mini-batch)
- ▶ The new datapoints are taken randomly from the batch

Stochastic gradient descent (SGD)

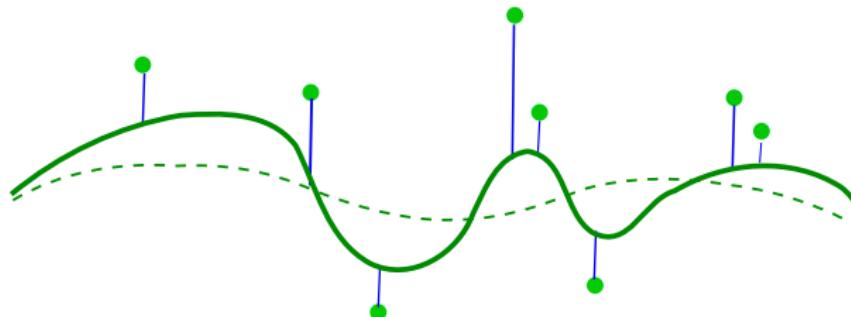


- ▶ Batch gradient descent is expensive when the batch is large
- ▶ We can decrease the cost by using several mini-batches (extra gain with parallel computations)
- ▶ In SGD, each mini-batch gives a different (inaccurate) estimate of the local gradient
- ▶ Combining several inaccurate estimates can help being accurate
- ▶ Note that SGD converges more slowly. Use it only for large batches.



Bottou, L. (2012) Stochastic gradient descent tricks. *Neural networks: Tricks of the trade*, pp. 421–436. Springer

Stochastic gradient descent (SGD)

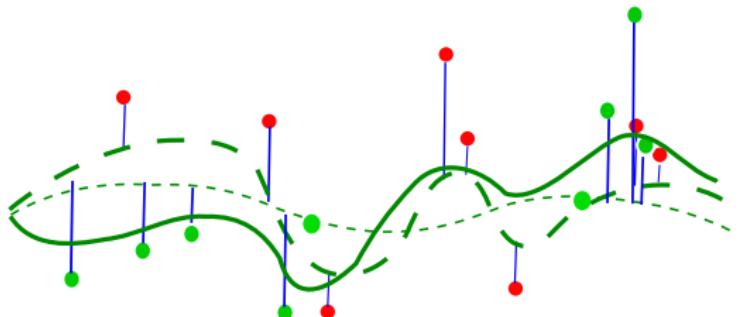


- ▶ Batch gradient descent is expensive when the batch is large
- ▶ We can decrease the cost by using several mini-batches (extra gain with parallel computations)
- ▶ In SGD, each mini-batch gives a different (inaccurate) estimate of the local gradient
- ▶ Combining several inaccurate estimates can help being accurate
- ▶ Note that SGD converges more slowly. Use it only for large batches.



Bottou, L. (2012) Stochastic gradient descent tricks. *Neural networks: Tricks of the trade*, pp. 421–436. Springer

Stochastic gradient descent (SGD)

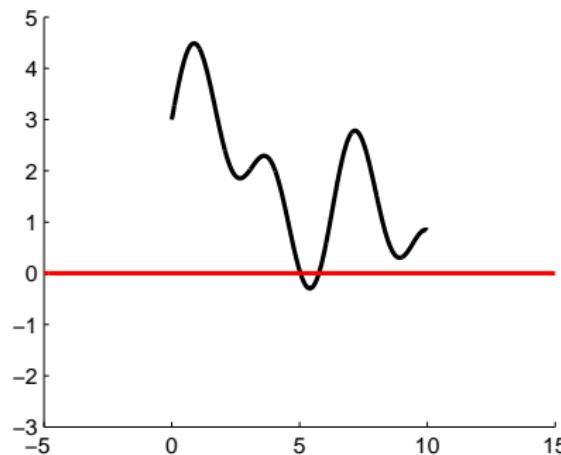


- ▶ Batch gradient descent is expensive when the batch is large
- ▶ We can decrease the cost by using several mini-batches (extra gain with parallel computations)
- ▶ In SGD, each mini-batch gives a different (inaccurate) estimate of the local gradient
- ▶ Combining several inaccurate estimates can help being accurate
- ▶ Note that SGD converges more slowly. Use it only for large batches.



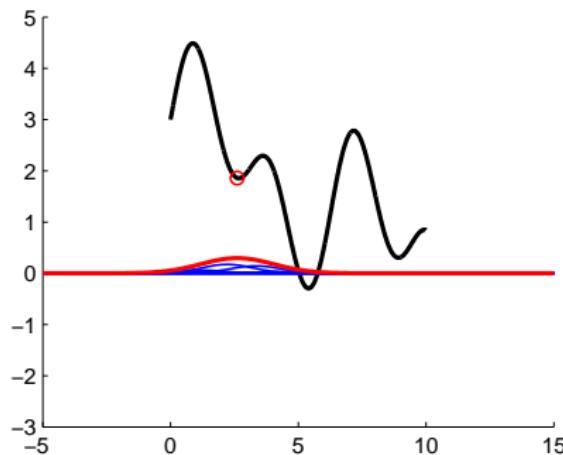
Bottou, L. (2012) Stochastic gradient descent tricks. *Neural networks: Tricks of the trade*, pp. 421–436. Springer

Stochastic gradient descent with RBFNs: Illustration



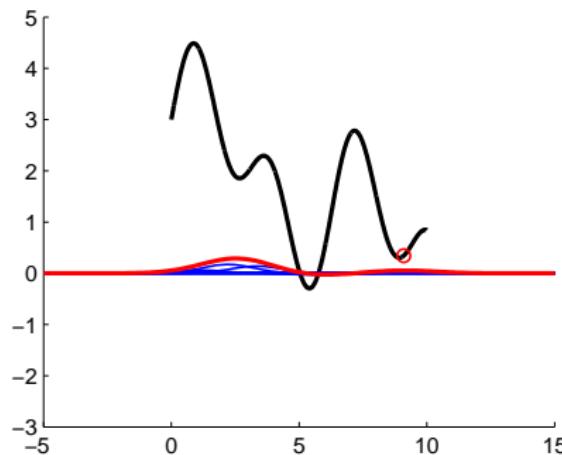
- Extreme case where we take one point of the batch at each iteration

Stochastic gradient descent with RBFNs: Illustration



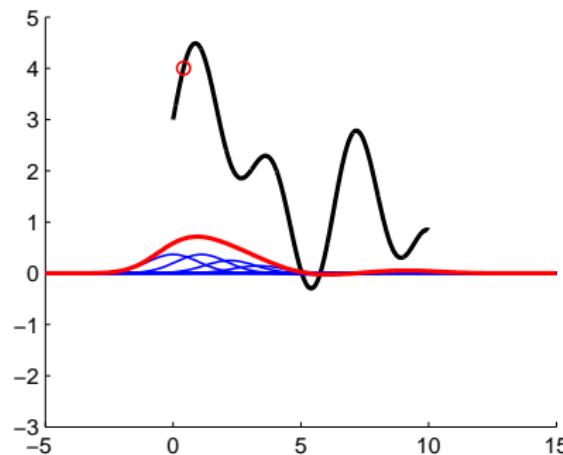
- Extreme case where we take one point of the batch at each iteration

Stochastic gradient descent with RBFNs: Illustration



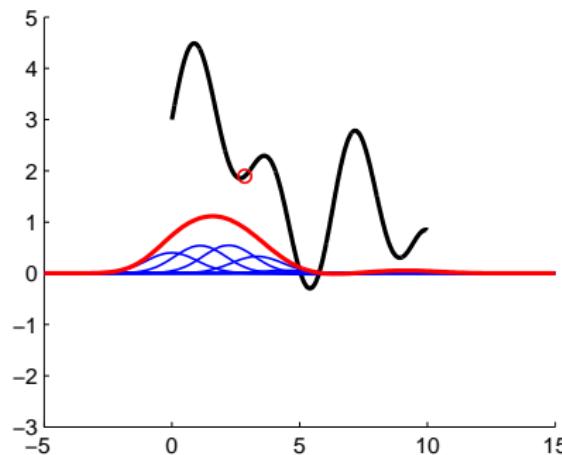
- Extreme case where we take one point of the batch at each iteration

Stochastic gradient descent with RBFNs: Illustration



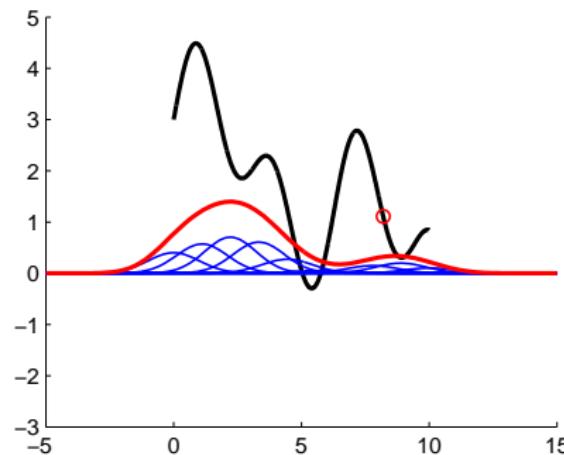
- Extreme case where we take one point of the batch at each iteration

Stochastic gradient descent with RBFNs: Illustration



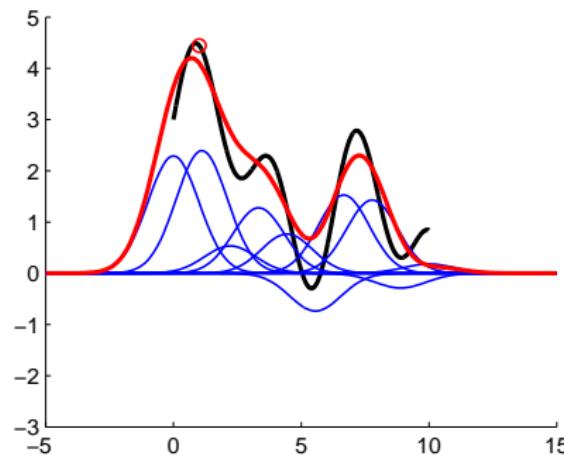
- Extreme case where we take one point of the batch at each iteration

Stochastic gradient descent with RBFNs: Illustration



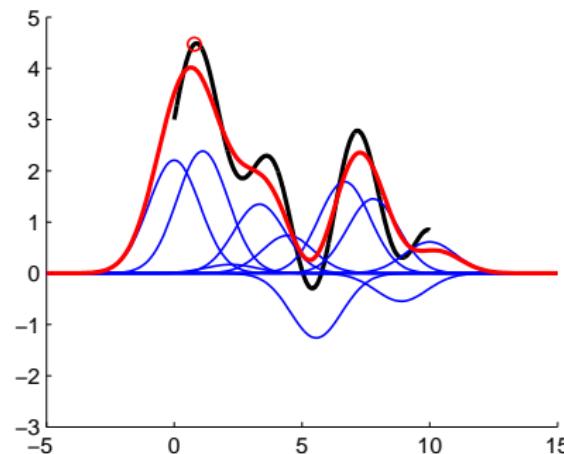
- Extreme case where we take one point of the batch at each iteration

Stochastic gradient descent with RBFNs: Illustration



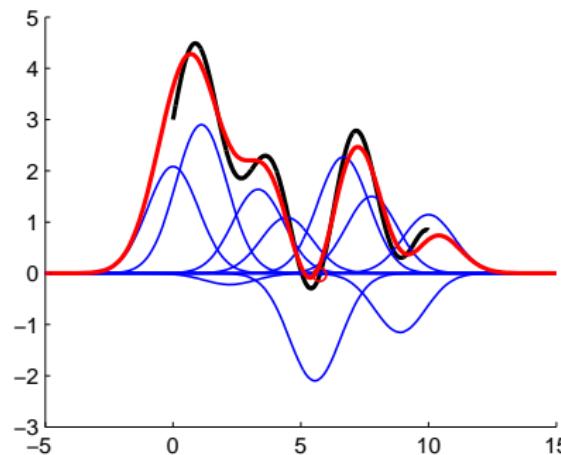
- Extreme case where we take one point of the batch at each iteration

Stochastic gradient descent with RBFNs: Illustration



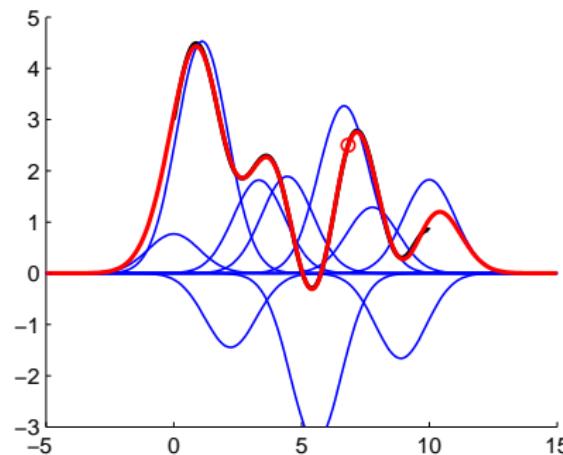
- Extreme case where we take one point of the batch at each iteration

Stochastic gradient descent with RBFNs: Illustration



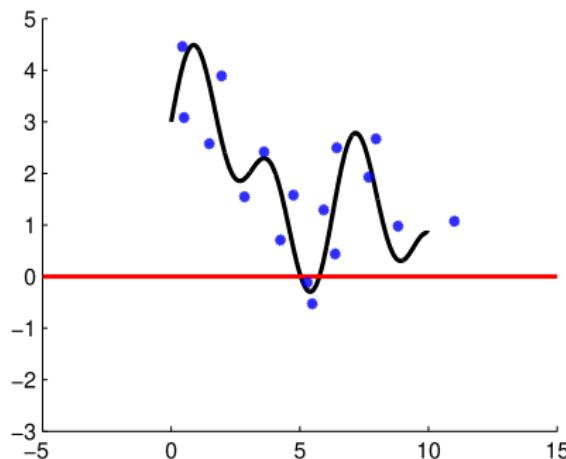
- Extreme case where we take one point of the batch at each iteration

Stochastic gradient descent with RBFNs: Illustration



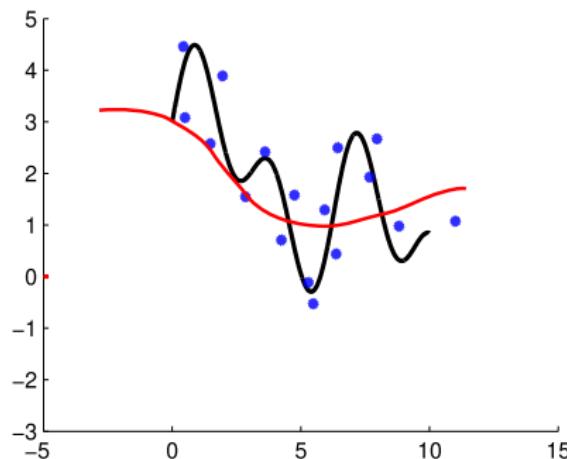
- Extreme case where we take one point of the batch at each iteration

Risk of overfitting



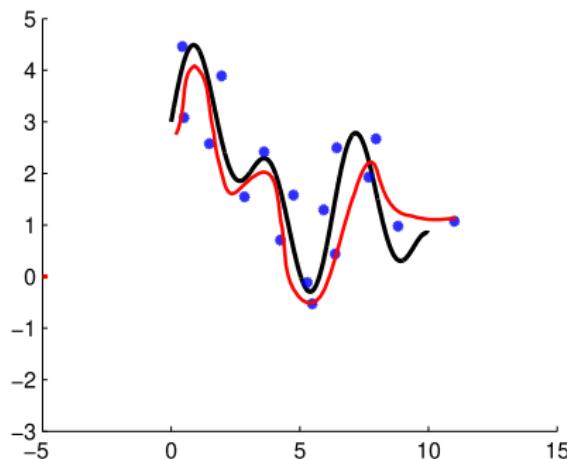
- ▶ In the previous simulation, measurements were noiseless
- ▶ What if we get noisy measurements?
- ▶ The risk is overfitting to data
- ▶ The question is when to stop

Risk of overfitting



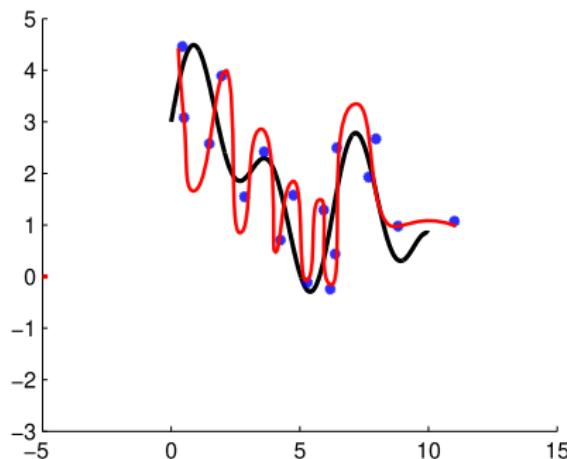
- ▶ In the previous simulation, measurements were noiseless
- ▶ What if we get noisy measurements?
- ▶ The risk is overfitting to data
- ▶ The question is when to stop

Risk of overfitting



- ▶ In the previous simulation, measurements were noiseless
- ▶ What if we get noisy measurements?
- ▶ The risk is overfitting to data
- ▶ The question is when to stop

Risk of overfitting



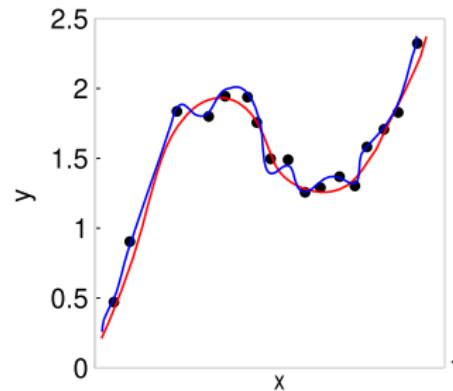
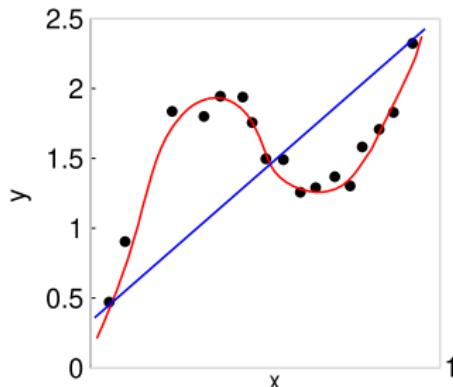
- ▶ In the previous simulation, measurements were noiseless
- ▶ What if we get noisy measurements?
- ▶ The risk is overfitting to data
- ▶ The question is when to stop

Fighting overfitting: Standard methodology



- ▶ Split the training set into training set and validation set
- ▶ Train on the training set
- ▶ Evaluate on the validation set
- ▶ When performance stops improving, stop iterating
- ▶ Should generalize well to the test set

Overfitting: influence of the model



- ▶ A model with more parameters is more prone to overfitting
- ▶ If the model does not have enough free parameters:
 - ▶ The training error may not go down towards 0
 - ▶ The validation performance may stop improving
- ▶ → Select a rich enough model...
- ▶ ...and use the validation trick to decide when to stop
- ▶ (Deep) neural networks are particularly rich models



Bottou, L. (2012).

Stochastic gradient descent tricks.

In *Neural networks: Tricks of the trade*, pages 421–436. Springer.



Schulman, J., Levine, S., Abbeel, P., Jordan, M. I., and Moritz, P. (2015).

Trust region policy optimization.

In Bach, F. R. and Blei, D. M., editors, *Proceedings of the 32nd International Conference on Machine Learning, ICML 2015, Lille, France, 6-11 July 2015*, volume 37 of *JMLR Workshop and Conference Proceedings*, pages 1889–1897. JMLR.org.